LTIC · UNIVERSITY OF BRITISH COLUMBIA

Working with AI tools, *at a glance*

A one-page rule of thumb for when it's OK to put information into AI assistants (Claude, Claude Code, ChatGPT, Copilot, and similar). When in doubt, err one category up and ask.

v1.0
Effective May 2026
Owner: LTIC
Review: annually

STOP

DO NOT SHARE

Don't put this into an AI tool

Anything that points to a real, identifiable person, or that UBC has a legal duty to protect (FIPPA) — including information that could be reconstructed to identify an individual. Models can also attach meaning to identifiers in ways that introduce bias, so strip names, IDs, and other identifiers from prompts unless they're absolutely necessary. If you wouldn't email it to a stranger, don't paste it into a chat with an unsupported service. Oversharing also creates IP reconstruction risk depending on the tool. **≈ ISS U1 High / Very High.**

- Student names, IDs, emails, grades
- Instructor personnel info, performance notes
- Health, financial, or HR information
- Passwords, API keys, tokens, secrets
- Identifiable course feedback or evaluations
- Accessibility / accommodation records
- Confidential meeting notes about people
- Unpublished research with identifiable subjects

!

SANITIZE FIRST

Weigh the harm — de-identify, or ask the owner

Internal-but-sensitive material **≈ ISS U1 Medium.** Before you share, weigh the potential harm and whether the tool actually needs identifying detail to do the work — de-identify wherever you can to limit bias and protect identity. Protected IP carries its own obligations: research data, Indigenous data (OCAP-governed), and other rights-bearing material may need permissions beyond FIPPA. If you're unsure how something is classified or what the risk looks like, ask the information owner before you paste. Use institutionally supported or endorsed tools where you can. Medium-risk work should run on an institutionally supported service, or get Administrative Head approval to execute a clickthrough agreement under Signing Resolution #26 (SR#26).

- De-identified or aggregated student data
- Draft policies, internal comms, planning docs
- Internal workflows & process docs
- Vendor contracts & procurement details
- Research data, Indigenous data (OCAP)
- Anonymized survey results
- Unpublished course materials & assessments
- Log / analytics data that could re-identify
- Anything marked "internal" or "draft"

GO

GENERALLY OK

Work freely — this is what AI is for

Anything that isn't about a specific person and doesn't contain secrets. You own the output — still review what comes back **≈ ISS U1 Low / public-equivalent.**

- UBC / LTIC source code & repositories
- Generic code, scripts, refactors (no secrets / keys)
- Generic templates, boilerplate, emails
- Literature review of public sources
- Learning a new tool, language, or concept
- Canvas API, Academic API, other UBC integrations
- Public-facing content already published
- Brainstorming, outlines, framing
- Documentation & technical writing
- System architecture, pseudocode, diagrams

Caveat — *green-zone safety assumes training is off on whatever tool you're using (see vendor cards below). With training on, institutional code raises copyright / IP questions — but unlike PII (red), it's never a FIPPA issue. Consumer Claude / Gemini / ChatGPT plans* **do not inherently meet UBC's ISS U9 contractual requirements***, so keep within ISS U1 Low on those tools.*

Helpful patterns WHEN IN DOUBT

1. **Ask AI to sanitize, locally.** Working with PII? Ask Claude to write a script that de-identifies the file on your machine — share only the cleaned output back.
2. **Aggregate before you ask.** Turn "students & grades" into "grade distribution by cohort." The pattern survives; the people don't.
3. **Keep secrets in `.env`.** Never paste keys, tokens, or DB strings into a chat — refer to them by name and let your code load them.
4. **Mind the border.** All major AI services process data outside Canada — that's a FIPPA issue for identifiable info. Sanitize first, or use an approved tool.
5. **Sandbox agentic tools.** Run Claude Code, Cursor, or similar in a DevContainer, Docker container, or WSL VM — isolate the agent's reach to a single project folder, not your whole filesystem.
6. **Least-privilege access.** Grant read-only tokens where possible. Work from a sanitized copy of the repo rather than a live branch. Scope service accounts to the task, not the organization.
7. **Exclude with ignore files.** Add a `.copilotignore` or `.cursorignore` (mirrors `.gitignore` syntax) to block secrets, config files, and credential stores from the agent's indexing context.
8. **Disable autorun.** Turn off autonomous execution modes in agentic tools. Maintain an allowlist of permitted shell commands — every action the agent proposes to run should require explicit human approval before it executes.

Before you deploy an agentic coding tool CSA AGENTIC TRUST FRAMEWORK

1. **Who are you?** Does the tool have a verified, auditable identity? Use a dedicated service account — not your personal credentials — so AI-generated activity is clearly attributable in audit logs.
2. **What are you doing?** Can you observe and log the agent's actions in real time? Agentic tools should produce an audit trail you can review. If you can't see what it did, you can't govern it.
3. **What are you eating?** What data is the agent ingesting — prompts, files, repo contents? Validate inputs before they reach the model. Sensitive files should be excluded via ignore rules before indexing begins.
4. **Where can you go?** What systems, files, and APIs can the agent reach? Enforce boundaries before you start. Sandboxing, read-only tokens, and scoped service accounts limit the blast radius of an error or a prompt injection.
5. **What if you go rogue?** Can you stop it? Know how to kill the session, revoke its credentials, and roll back any changes it made. Disable autorun so you retain a human-approval gate before anything executes.

Based on the CSA Agentic Trust Framework (Feb 2026) — Identity, Behavior, Data Governance, Segmentation, Incident Response.

AI tools at a glance WHAT THE POLICIES ACTUALLY SAY

verified May 2026

Claude · Claude Code FREE · PRO · MAX · TEAM · ENTERPRISE · API

Default — training off
Not used for training unless you opt in at `…/data-privacy-controls` — same toggle for Free, Pro, Max (not a paid/free split). Team / Enterprise / API never trained on. **Always validate that "Help improve Claude" is disabled.** Retention: **30 days** default, **5 years** if opted in. **Keeping chat history ≠ training** — they're separate settings. Safety-flagged chats reviewed regardless; data processed in the **US**.

⚠ Claude Code (CLI)
Stores session transcripts in **plaintext** at `~/.claude/projects/` for 30 days — anything you paste lives there too. Tune with `cleanupPeriodDays`. Running `/feedback` uploads a transcript (5y retention) **and is used for training regardless of your opt-out**; disable with `DISABLE_FEEDBACK_COMMAND=1`.

⚠ claude.ai (web / app)
Memory (all plans): Settings → Capabilities → "*Generate memory from chat history*". Pause, reset, or view what's stored. Memory inherits chat retention; Anthropic's docs don't explicitly address whether memory entries themselves train models — assume in-scope if your training toggle is on. **Thumbs-up/down feedback** on a reply uses that conversation for training regardless of opt-out. **Incognito chats** (ghost icon) skip both memory and training regardless of settings — use for ad-hoc sensitive questions.

ISS U1 ceiling on this plan — Low. *Medium use needs Anthropic's commercial / Enterprise terms + an Assessment of Use Case + ISS U9 / SR#26 review where necessary.*

Sources — Is my data used for model training? (Mar 2026) · Claude Code: Data Usage · Anthropic Privacy Policy (Jan 2026) · Anthropic Commercial Terms (Jun 2025)

Gemini · Gemini CLI APPS · CLI · WORKSPACE · VERTEX

Default — varies by sign-in
Google does not provide a single global toggle. Whether your prompts are used for training depends entirely on **which account or API key you authenticated with**. The two columns to the right break it down.

⚠ Gemini Apps (web / mobile)
Personal accounts (default): with **Keep Activity ON**, chats may be used to train models and a subset is reviewed by humans — toggle off in account settings to stop both. **Workspace / school accounts:** training off by default, no human review without permission. Reviewed chats kept up to **3 years**; **72 hours** minimum either way.

⚠ Gemini CLI — auth method matters
Trained on by default: personal Google account; free Gemini API key. **Personal account:** disabling the CLI's **Usage Statistics** setting stops prompt / code collection (it covers both telemetry and training data). **Free API key:** training continues regardless — the Gemini API Additional Terms govern data sent to the API endpoint and have no opt-out. **Not trained on:** Workspace / Enterprise account, paid Gemini API, Vertex AI.

ISS U1 ceiling on this plan — Low. *Medium use needs Workspace / Enterprise auth or Vertex AI + an Assessment of Use Case + ISS U9 / SR#26 review where necessary.*

Sources — Gemini Apps Privacy Hub (May 2026) · Gemini CLI ToS & Privacy · Gen-AI in Workspace Privacy Hub · Gemini for Google Cloud — Data Governance

ChatGPT · Codex FREE · PLUS · PRO · TEAM · ENTERPRISE · EDU · API

⚠ Default depends on plan
Free, Plus, Pro: training is **ON by default**. Turn it off at Settings → Data Controls → *Improve the model for everyone*. **Team, Enterprise, Edu, API:** training is **OFF by default** (no toggle required). **Caveat:** submitting thumbs-up/down feedback on a chat opts that entire conversation into training regardless of the toggle.

Temporary Chat
A per-conversation mode that's **not used for training** and **deleted after 30 days**. Useful for one-off questions on a consumer account when you can't easily switch tools — still sanitize identifiable info first.

⚠ Memory feature
Persists context (your facts, preferences) **across all chats. Memory and training are independent toggles** — you can disable memory while training stays on, or vice versa. Memory controls: Settings → Personalization → Memory (clear, disable saving, disable reference). Memorized content is in scope for training only if "Improve the model for everyone" is on.

ISS U1 ceiling on this plan — Low. *Medium use needs Team / Enterprise / Edu + an Assessment of Use Case + ISS U9 / SR#26 review where necessary.*

Sources — Data Controls FAQ · OpenAI Privacy Policy · OpenAI Enterprise Privacy · Keep history on, disable training

Rule of thumb — *If it identifies a person or can't be shared publicly, don't put it in a consumer AI tool.*

Questions? Ask in #ltic-genai-code · LT.hub@ubc.ca
UBC ISS · cio.ubc.ca/information-security-standards
UBC GenAI Principles · genai.ubc.ca/guidance
PIA Guidelines · privacymatters.ubc.ca/pia-stra/pia-guidelines-tools