

# Working with AI tools, *at a glance*

A one-page rule of thumb for when it's OK to put information into AI assistants (Claude, Claude Code, ChatGPT, Copilot, and similar). When in doubt, err one category up and ask.

v1.1

Effective August 2026

Owner: LTIC

Review: annually



**DO NOT  
SHARE**

## Don't put this into an AI tool

Anything that points to a real, identifiable person, or that UBC has a legal duty to protect (FIPPA) — including information that could be reconstructed to identify an individual. Models can also attach meaning to identifiers in ways that introduce bias, so strip names, IDs, and other identifiers from prompts unless they're absolutely necessary. If you wouldn't email it to a stranger, don't paste it into a chat with an unsupported service. Oversharing also creates IP reconstruction risk depending on the tool. **≈ ISS U1 High / Very High.**

- Student names, IDs, emails, grades
- Instructor personnel info, performance notes
- Health, financial, or HR information
- Passwords, API keys, tokens, secrets
- Identifiable course feedback or evaluations
- Accessibility / accommodation records
- Confidential meeting notes about people
- Unpublished research with identifiable subjects



**SANITIZE  
FIRST**

## Weigh the harm — de-identify, or ask the owner

Internal-but-sensitive material **≈ ISS U1 Medium**. Before you share, weigh the potential harm and whether the tool actually needs identifying detail to do the work — de-identify wherever you can to limit bias and protect identity. Protected IP carries its own obligations: research data, Indigenous data (OCAP-governed), and other rights-bearing material may need permissions beyond FIPPA. If you're unsure how something is classified or what the risk looks like, ask the information owner before you paste. Use institutionally supported or endorsed tools where you can. Medium-risk work should run on an institutionally supported service, or get Administrative Head approval to execute a clickthrough agreement under Signing Resolution #26 (SR#26).

- ◆ De-identified or aggregated student data
- ◆ Draft policies, internal comms, planning docs
- ◆ Internal workflows & process docs
- ◆ Vendor contracts & procurement details
- ◆ Research data, Indigenous data (OCAP)
- ◆ Anonymized survey results
- ◆ Unpublished course materials & assessments
- ◆ Log / analytics data that could re-identify
- ◆ Anything marked "internal" or "draft"



**GENERALLY  
OK**

## Work freely — this is what AI is for

Anything that isn't about a specific person and doesn't contain secrets. You own the output — still review what comes back **≈ ISS U1 Low / public-equivalent**.

- UBC / LTIC source code & repositories
- Generic code, scripts, refactors (no secrets / keys)
- Generic templates, boilerplate, emails
- Literature review of public sources
- Learning a new tool, language, or concept
- Canvas API, Academic API, other UBC integrations
- Public-facing content already published
- Brainstorming, outlines, framing
- Documentation & technical writing
- System architecture, pseudocode, diagrams

**Caveat** — green-zone safety assumes training is off on whatever tool you're using (see vendor cards below). With training on, institutional code raises copyright / IP questions — but unlike PII (red), it's never a FIPPA issue. Consumer Claude / Gemini / ChatGPT plans **do not inherently meet UBC's ISS U9 contractual requirements**, so keep within ISS U1 Low on those tools.

## Helpful patterns WHEN IN DOUBT

- 1. Ask AI to sanitize, locally.** Working with PII? Ask Claude to write a script that de-identifies the file on your machine — share only the cleaned output back.
- 2. Aggregate before you ask.** Turn "students & grades" into "grade distribution by cohort." The pattern survives; the people don't.
- 3. Keep secrets in .env.** Never paste keys, tokens, or DB strings into a chat — refer to them by name and let your code load them.
- 4. Mind the border.** Consumer AI tiers process data outside Canada — a FIPPA issue for identifiable info. Some enterprise tiers offer Canadian at-rest residency, but inference still runs in the US. Sanitize first, or use an approved tool.
- 5. Sandbox agentic tools.** Run Claude Code, Cursor, or similar in a DevContainer, Docker container, or WSL VM — isolate the agent's reach to a single project folder, not your whole filesystem.
- 6. Least-privilege access.** Grant read-only tokens where possible. Work from a sanitized copy of the repo rather than a live branch. Scope service accounts to the task, not the organization.
- 7. Exclude sensitive paths.** Cursor: a `.cursorignore` (`.gitignore` syntax, best-effort). Copilot: content exclusion is set in GitHub repo / org settings (paid plans) — and Copilot's agent modes don't honor it, so don't rely on exclusions as your only guard.
- 8. Disable autorun.** Turn off autonomous execution modes in agentic tools. Maintain an allowlist of permitted shell commands — every action the agent proposes to run should require explicit human approval before it executes.

## Before you deploy an agentic coding tool CSA AGENTIC TRUST FRAMEWORK

1. **Who are you?** Does the tool have a verified, auditable identity? Use a dedicated service account — not your personal credentials — so AI-generated activity is clearly attributable in audit logs.
2. **What are you doing?** Can you observe and log the agent's actions in real time? Agentic tools should produce an audit trail you can review. If you can't see what it did, you can't govern it.
3. **What are you eating?** What data is the agent ingesting — prompts, files, repo contents? Validate inputs before they reach the model. Sensitive files should be excluded via ignore rules before indexing begins.
4. **Where can you go?** What systems, files, and APIs can the agent reach? Enforce boundaries before you start. Sandboxing, read-only tokens, and scoped service accounts limit the blast radius of an error or a prompt injection.
5. **What if you go rogue?** Can you stop it? Know how to kill the session, revoke its credentials, and roll back any changes it made. Disable autorun so you retain a human-approval gate before anything executes.

Based on the [CSA Agentic Trust Framework](#) (Feb 2026) — Identity, Behavior, Data Governance, Segmentation, Incident Response.

## AI tools at a glance WHAT THE POLICIES ACTUALLY SAY

verified August 2026

### Claude • Claude Code FREE • PRO • MAX • TEAM • ENTERPRISE • API

#### Default — training off

Not used for training unless you opt in at `.../data-privacy-controls` — same toggle for Free, Pro, Max (not a paid/free split). Team / Enterprise / API never trained on. **Always validate that "Help improve Claude" is disabled.** Retention: **30 days** default, **5 years** if opted in. **Keeping chat history ≠ training** — they're separate settings. Safety-flagged chats may be reviewed and used for training regardless; data processed in the US.

#### △ Claude Code (CLI)

Stores session transcripts in **plaintext** at `~/claude/projects/` for 30 days — anything you paste lives there too. Tune with `cleanupPeriodDays`. Running `/feedback` (also `/bug`, `/share`) uploads a transcript (5y retention) **usable for training regardless of your opt-out**; disable with `DISABLE_FEEDBACK_COMMAND=1`. The "How is Claude doing?" survey's optional transcript share is kept 6 months but never trains; silence it with `CLAUDE_CODE_DISABLE_FEEDBACK_SURVEY=1`.

#### △ claude.ai (web / app)

**Memory** (all plans, **on by default**): Settings → Capabilities → "Generate memory from chat history". Pause, reset, or view what's stored. Memory inherits chat retention; Anthropic's docs don't explicitly address whether memory entries themselves train models — assume in-scope if your training toggle is on. **Thumbs-up/down feedback** on a reply uses that conversation for training regardless of opt-out. **Incognito chats** (ghost icon) skip both memory and training regardless of settings — use for ad-hoc sensitive questions.

ISS U1 ceiling on this plan — **Low**. Medium use needs Anthropic's commercial / Enterprise terms + an Assessment of Use Case + ISS U9 / SR#26 review where necessary.

Sources — [Is my data used for model training?](#) (Mar 2026) · [Claude Code: Data Usage](#) · [Anthropic Privacy Policy](#) (Jun 2026) · [Anthropic Commercial Terms](#) (Jun 2025)

### Gemini • Gemini CLI • Antigravity APPS • CLI • ANTIGRAVITY • WORKSPACE • VERTEX

#### Default — varies by sign-in

Google does not provide a single global toggle. Whether your prompts are used for training depends entirely on **which account or API key you authenticated with**. The two columns to the right break it down.

#### △ Gemini Apps (web / mobile)

**Personal accounts (default):** with **Keep Activity ON**, chats may be used to train models and a subset is reviewed by humans — toggle off in account settings to stop training use (safety / abuse review can still occur). **Workspace / school accounts:** training off by default, no human review without permission. Reviewed chats kept up to **3 years; 72 hours** minimum either way.

#### △ Gemini CLI & Antigravity — auth matters

**Gemini CLI / Code Assist:** requires Workspace / Enterprise auth, a paid Gemini API key, or Vertex AI — **none trained on**. There is no personal-account sign-in. **Free Gemini API key:** trained on with **no opt-out** — the Gemini API Additional Terms have no off switch. **Antigravity** (Google's consumer agentic IDE): individual accounts are trained on **by default** — opt out in its settings; it can also route requests to Anthropic models under separate terms.

ISS U1 ceiling on this plan — **Low**. Medium use needs Workspace / Enterprise auth or Vertex AI + an Assessment of Use Case + ISS U9 / SR#26 review where necessary.

Sources — [Gemini Apps Privacy Hub](#) (Aug 2026) · [Gemini CLI ToS & Privacy](#) · [Antigravity Terms](#) · [Gen-AI in Workspace Privacy Hub](#) · [Gemini for Google Cloud — Data Governance](#)

### ChatGPT • Codex FREE • GO • PLUS • PRO • BUSINESS • ENTERPRISE • EDU • API

#### △ Default depends on plan

**Free, Go, Plus, Pro:** training is **ON by default**. Turn it off at Settings → Data Controls → *Improve the model for everyone* (the toggle also covers Codex tasks; voice mode has its own). **Business, Enterprise, Edu, API:** training is **OFF by default** (no toggle required). **Caveats:** thumbs-up/down feedback makes that entire conversation usable for training regardless of the toggle; Free / Go tiers show **ads** targeted using your past chats (separate ad-personalization toggle).

#### Temporary Chat

A per-conversation mode that's **not used for training** and **deleted after 30 days**. Useful for one-off questions on a consumer account when you can't easily switch tools — still sanitize identifiable info first.

#### △ Memory feature

Auto-builds a **memory summary** from your chats, files, and connected apps, referenced **across all chats**. **Memory and training are independent toggles**. Controls: Settings → Personalization → Memory (clear, disable saving, disable reference). Deleting memories doesn't delete the source chats — memory can rebuild from history if re-enabled. Memorized content is in scope for training only if "Improve the model for everyone" is on.

ISS U1 ceiling on this plan — **Low**. Medium use needs Business / Enterprise / Edu + an Assessment of Use Case + ISS U9 / SR#26 review where necessary.

Sources — [Data Controls FAQ](#) · [OpenAI Privacy Policy](#) · [OpenAI Enterprise Privacy](#) · [Keep history on, disable training](#)

**Rule of thumb** — *If it identifies a person or can't be shared publicly, don't put it in a consumer AI tool.*

Questions? Ask in [#ltic-genai-code](#) · [LT.hub@ubc.ca](mailto:LT.hub@ubc.ca)  
UBC ISS · [cio.ubc.ca/information-security-standards](https://cio.ubc.ca/information-security-standards)  
UBC GenAI Principles · [genai.ubc.ca/guidance](https://genai.ubc.ca/guidance)  
PIA Guidelines · [privacymatters.ubc.ca/pia-stra/pia-guidelines-tools](https://privacymatters.ubc.ca/pia-stra/pia-guidelines-tools)